

Methods

Methods I

The template/ architecture is deliberately bifurcated into a globally shared infrastructure/ layer and project-specific projects/ silos. This section describes the four core design patterns, the YAML-declared pipeline (16 stages; default full 10) pipeline that operationalizes them, and the AI collaboration model that distinguishes this system from conventional research templates.

The Two-Layer Architecture I

The repository is organized into two strictly separated layers:

Infrastructure Layer (infrastructure/): 28 infrastructure subdirectories—25 of them independently-importable Python packages—comprising ~708 modules and providing reusable services. Each importable package has its own `__init__.py`, `AGENTS.md`, and `README.md`, and exports a well-defined public API (the remaining subdirectories, e.g. `config/`, hold shared configuration). The infrastructure layer knows nothing about any specific project—it provides generic capabilities (logging, rendering, validation, steganography) that any project may consume.

Project Layer (projects/): Self-contained research workspaces. Each project directory contains:

Directory	Purpose
manuscript/	Markdown chapters and <code>config.yaml</code>
scripts/	Thin orchestrator scripts (Stage 02)
src/	Project-specific Python modules
tests/	Project-specific test suite

The Two-Layer Architecture II

data/	Input datasets and generated data
output/	Pipeline artifacts: PDF, figures, reports, logs
docs/	Project-specific architecture documentation

The Two-Layer Architecture III

The two layers communicate exclusively through Python imports and filesystem paths. No project modifies infrastructure code; no infrastructure module references a specific project by name (except via runtime project discovery).

The Standalone Project Paradigm I

Projects are designed to be completely self-contained. Adding a new project requires no changes to the infrastructure layer, no modifications to `pyproject.toml`, and no updates to the pipeline orchestrator. A project is automatically discovered if and only if it satisfies two conditions:

1. It exists as a subdirectory of `projects/`.
2. It contains the file `manuscript/config.yaml`.

The Standalone Project Paradigm II

This paradigm enables horizontal scaling: N researchers can maintain N independent projects within a single repository, sharing infrastructure without coupling. Each project declares its own testing tolerances, manuscript metadata, LLM review preferences, and rendering configuration in its `config.yaml`. The system currently hosts its public canonical exemplars under `projects/templates/` (`templates/template_active_inference`, `templates/template_advanced_literature_review`, `templates/template_autopoiesis`, `templates/template_autoresearch_project`, `templates/template_autoscientists`, `templates/template_code_project`, `templates/template_data_descriptor`, `templates/template_eda_notebook`, `templates/template_formal`, `templates/template_gold_refinement`, `templates/template_literature_meta_analysis`, `templates/template_madlib`, `templates/template_methods_paper`, `templates/template_newspaper`, `templates/template_pitch_deck`, `templates/template_pools_rules_tools`, `templates/template_prose_project`, `templates/template_redacted_report`, `templates/template_registered_report`, `templates/template_search_pro`

The Thin Orchestrator Pattern I

All scripts in `scripts/` (both infrastructure-level and project-level) follow the Thin Orchestrator pattern [gamma1995design]:

- ▶ **No domain logic:** Scripts contain zero algorithmic implementation. They import functions from `src/` modules and wire them to infrastructure services.
- ▶ **Configuration-driven:** Behavior is parameterized by `config.yaml`, not by hardcoded values.
- ▶ **Stateless:** Scripts read inputs, call functions, write outputs. They maintain no persistent state between invocations.
- ▶ **Logged:** Every significant action is logged via `infrastructure.core.logging.utils.get_logger`.

The Thin Orchestrator Pattern II

This pattern ensures that all testable logic lives in `src/` where it is subject to the Zero-Mock testing policy, while scripts remain thin enough to be audited by visual inspection. The separation draws on the Mediator pattern from Gamma et al. [[@gamma1995design](#)], where scripts mediate between infrastructure services and project-specific code without implementing any logic of their own.

To make this concrete, the following contrasts the anti-pattern with the correct pattern:

```
# ANTI-PATTERN: domain logic embedded in script  
def calculate_average(data): # ← never put computation in script  
    return sum(data) / len(data)
```

```
result = calculate_average([1, 2, 3])
```

```
# CORRECT PATTERN: script imports from src/ and only wires  
from projects.my_project.src.statistics import calculate_average
```

The Thin Orchestrator Pattern III

```
result = calculate_average([1, 2, 3])  # ← scripts wire, n
```

The critical property is that `calculate_average` in the correct pattern lives in a testable `src/` module, is covered by the Zero-Mock test suite, and can be independently imported, tested, and reused—whereas the anti-pattern buries logic in a script that is invisible to coverage tools.